

Carnegie Mellon University

ADVANCED DATABASE SYSTEMS

OLTP Indexes
(Whole-Key Data Structures)

@Andy_Pavlo // 15-721 // Spring 2020

UPCOMING DATABASE EVENTS

Snowflake Optimizer Talk

- Monday Feb 3rd @ 4:30pm
- GHC 9115



WHOLE-KEY DATA STRUCTURE

A "whole-key" order preserving data structure stores all the digits of a key together in nodes.

→ A worker thread has to compare the entire search key with keys in the data structure during traversal.

We will discuss "partial-key" data structures (i.e., tries) next class.

TODAY'S AGENDA

In-Memory T-Tree

Latch-Free Bw-Tree

B+Tree Optimistic Latching



OBSERVATION

The original B+Tree was designed for efficient access of data stored on slow disks.

Is there an alternative data structure that is specifically designed for in-memory databases?



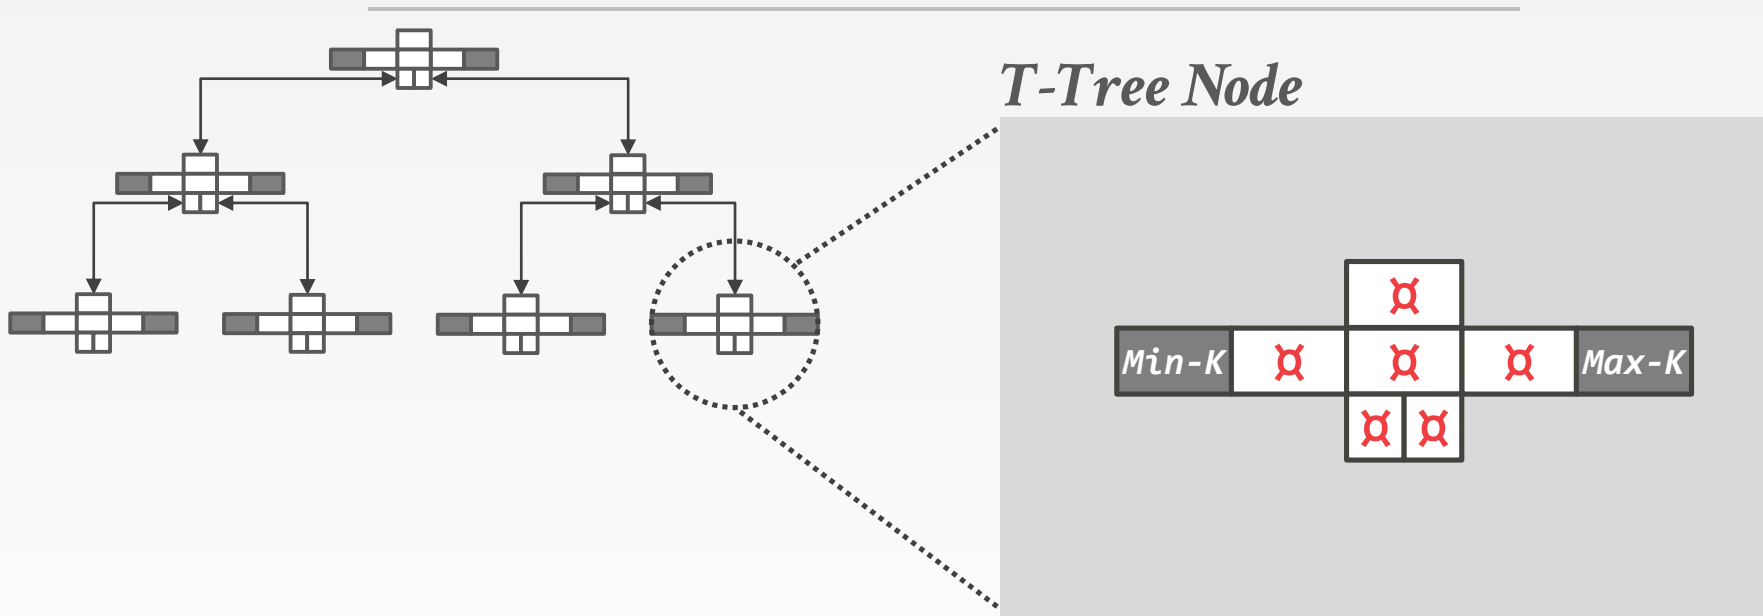
T-TREES

Based on AVL Trees. Instead of storing keys in nodes, store pointers to their original values.

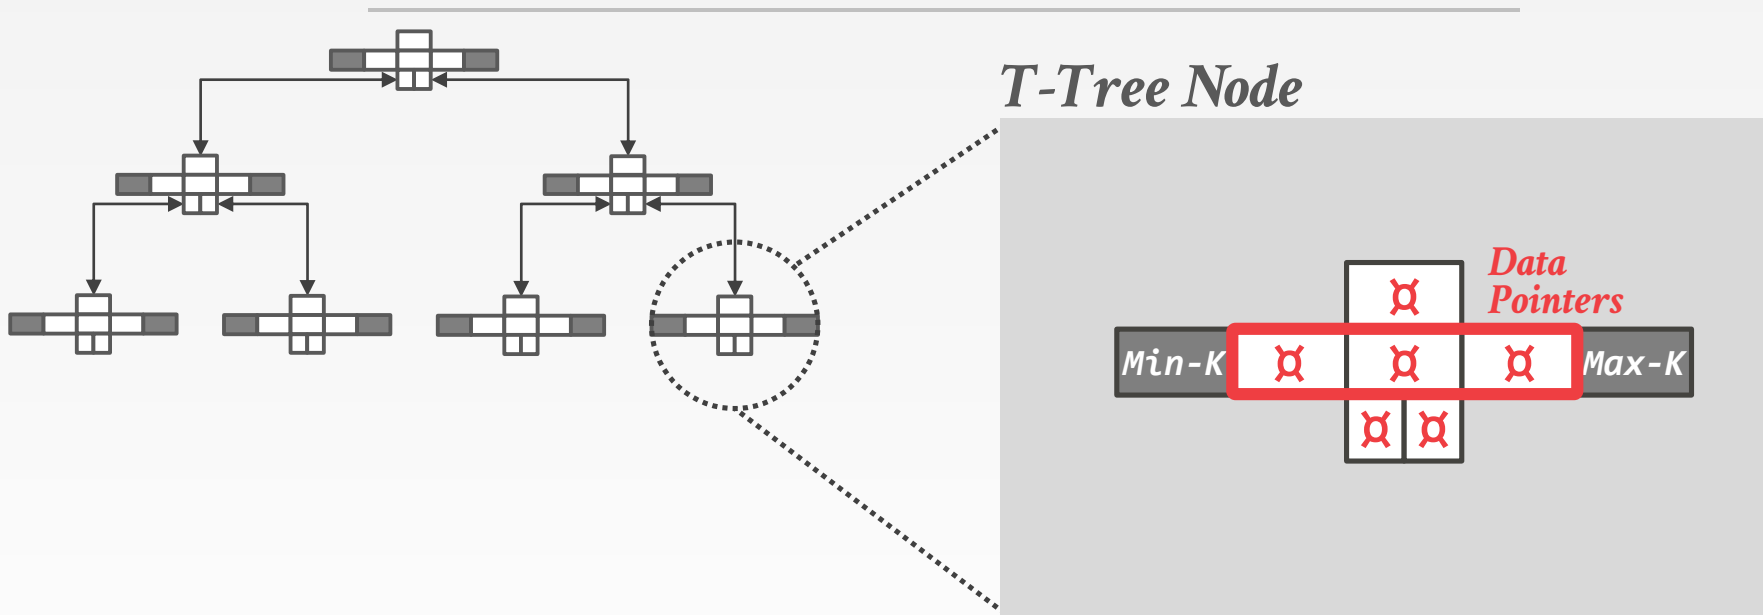
Proposed in 1986 from Univ. of Wisconsin
Used in TimesTen and other early in-memory DBMSs during the 1990s.



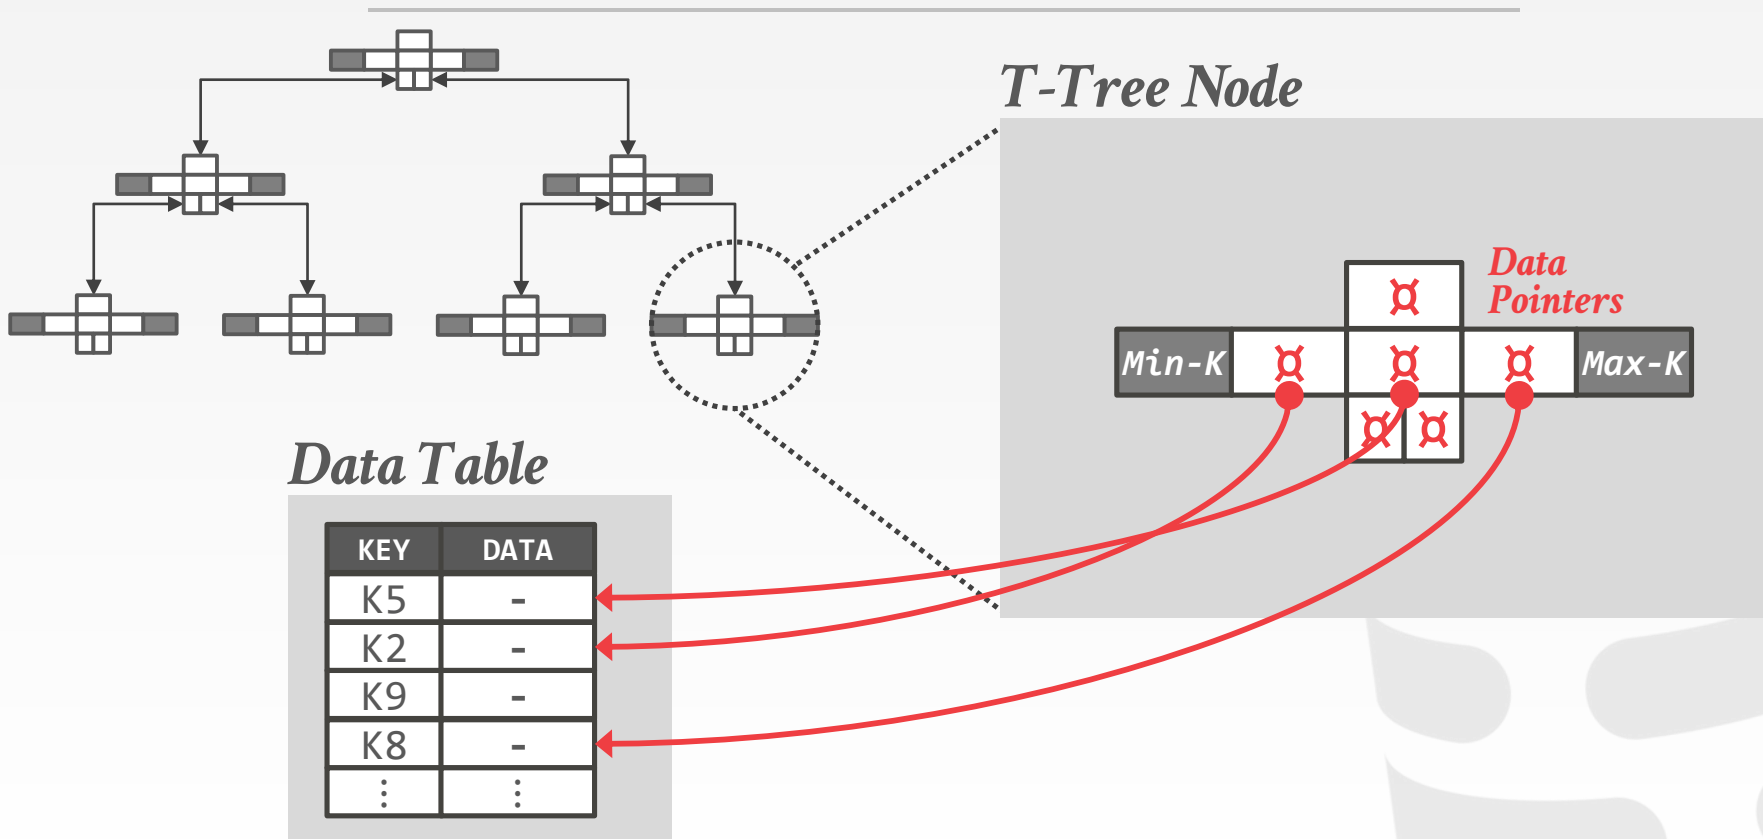
T-TREES




T-TREES



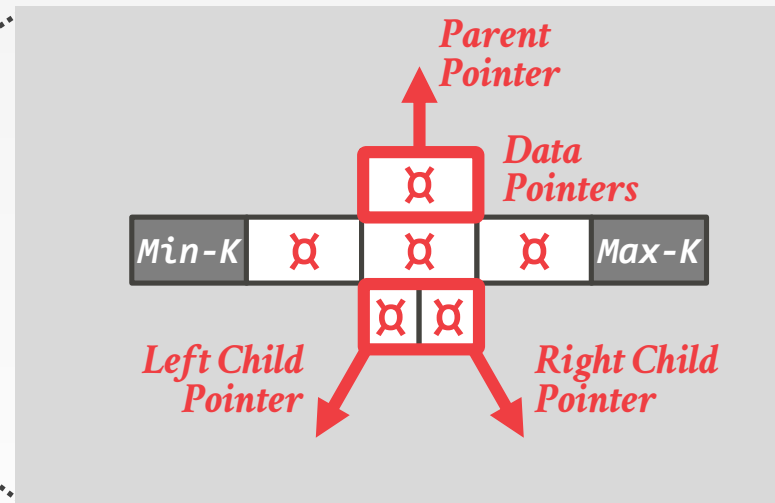
T-TREES



T-TREES




T-Tree Node



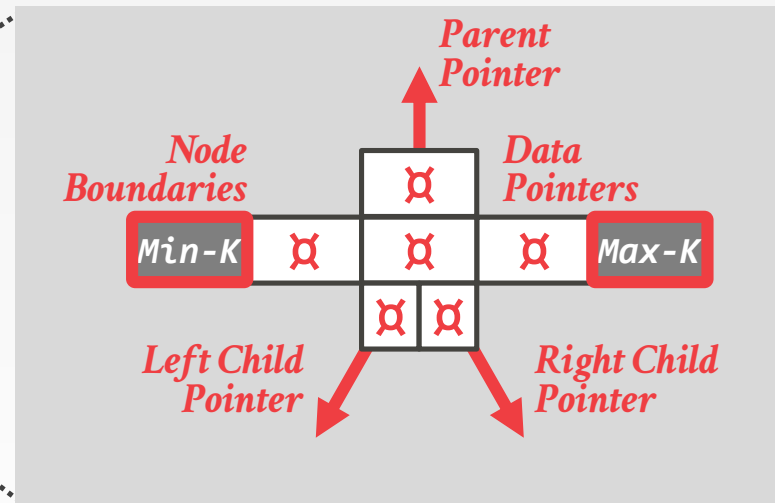
Data Table

KEY	DATA
K5	-
K2	-
K9	-
K8	-
⋮	⋮

T-TREES



T-Tree Node




Data Table

KEY	DATA
K5	-
K2	-
K9	-
K8	-
⋮	⋮

T-TREES: SEARCH

Find K2




Data Table

KEY	DATA
K1	-
K2	-
K3	-
K4	-
K5	-
K6	-
K7	-
K8	-
K9	-

T-TREES: SEARCH

Find K2




Data Table

KEY	DATA
K1	-
K2	-
K3	-
K4	-
K5	-
K6	-
K7	-
K8	-
K9	-

T-TREES: SEARCH

Find K2




Data Table

KEY	DATA
K1	-
K2	-
K3	-
K4	-
K5	-
K6	-
K7	-
K8	-
K9	-

T-TREES: SEARCH

Find K2




Data Table

KEY	DATA
K1	-
K2	-
K3	-
K4	-
K5	-
K6	-
K7	-
K8	-
K9	-

T-TREES: SEARCH

Find K2




Data Table

KEY	DATA
K1	-
K2	-
K3	-
K4	-
K5	-
K6	-
K7	-
K8	-
K9	-


T-TREES: SEARCH

Find K2



T-TREES: SEARCH

Find K2



T-TREE: ADVANTAGES

Uses less memory because it does not store keys inside of each node.

The DBMS evaluates all predicates on a table at the same time when accessing a tuple (i.e., not just the predicates on indexed attributes).

T-TREES: DISADVANTAGES

Difficult to rebalance.

Difficult to implement safe concurrent access.

Must chase pointers when scanning range or performing binary search inside of a node.

→ This greatly hurts cache locality.

OBSERVATION

Because CaS only updates a single address at a time, this limits the design of our data structures
→ We cannot build a latch-free B+Tree because we need to update multiple pointers on split/merge operations.

What if we had an indirection layer that allowed us to update multiple addresses atomically?

BW-TREE

Latch-free B+Tree index built for the Microsoft Hekaton project.

Key Idea #1: Deltas

- No updates in place
- Reduces cache invalidation.

Key Idea #2: Mapping Table

- Allows for CaS of physical locations of pages.



THE BW-TREE: A B-TREE FOR NEW HARDWARE
ICDE 2013


BW-TREE: MAPPING TABLE

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

*Logical
Pointer* 


*Physical
Pointer* 



BW-TREE: MAPPING TABLE

Mapping Table

PID	Addr
101	●
102	●
103	
104	●



BW-TREE: DELTA UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

▲ Insert K0

Page 102

*Logical
Pointer* →

*Physical
Pointer* →

Each update to a page produces a new delta.


BW-TREE: DELTA UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

*Logical
Pointer* 

*Physical
Pointer* 

 **Insert K0**

Page 102

Each update to a page produces a new delta.


Delta physically points to base page.


Install delta address in physical address slot of mapping table using CaS.


BW-TREE: DELTA UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	

*Logical
Pointer* 

*Physical
Pointer* 

 **Insert K0**

Page 102

Each update to a page produces a new delta.

Delta physically points to base page.

Install delta address in physical address slot of mapping table using CaS.


BW-TREE: DELTA UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ----->



Each update to a page produces a new delta.

Delta physically points to base page.

Install delta address in physical address slot of mapping table using CaS.


BW-TREE: DELTA UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>



Each update to a page produces a new delta.

Delta physically points to base page.

Install delta address in physical address slot of mapping table using CaS.


BW-TREE: SEARCH

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

*Logical
Pointer* ----->

*Physical
Pointer* ----->



Traverse tree like a regular B+tree.

If mapping table points to delta chain, stop at first occurrence of search key.

Otherwise, perform binary search on base page.

BW-TREE: CONTENTION UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

▲ Insert K0

Page 102

*Logical
Pointer* →

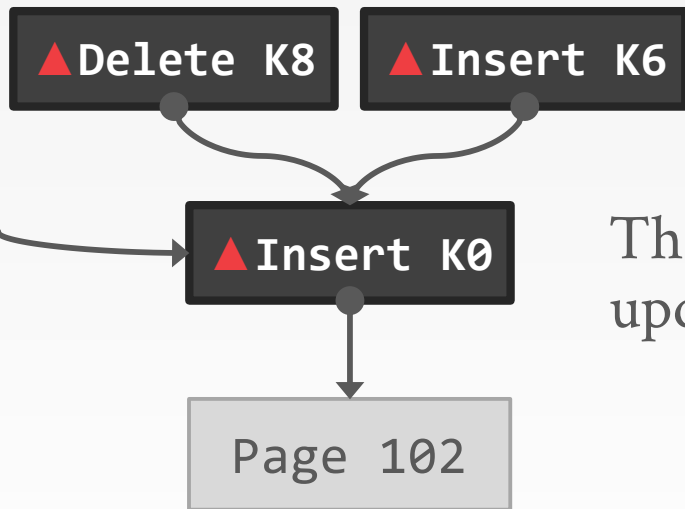
*Physical
Pointer* →

Threads may try to install updates to same page.

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	●
103	
104	



Threads may try to install updates to same page.

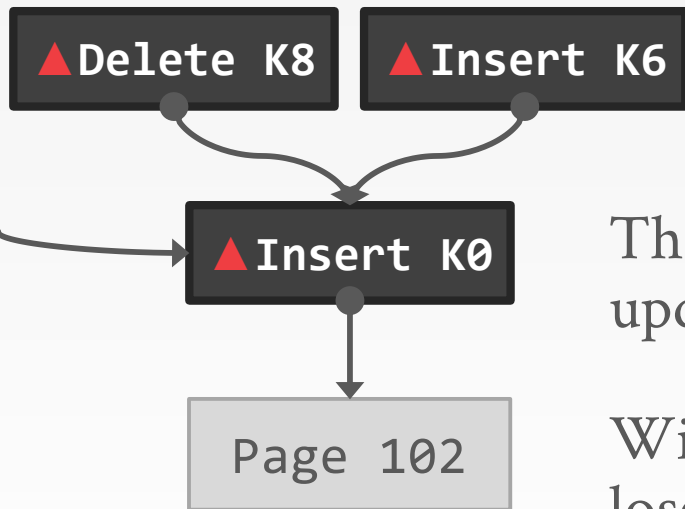
Logical
Pointer ----->

Physical
Pointer ————>

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	●
103	
104	



Threads may try to install updates to same page.

Winner succeeds, any losers must retry or abort

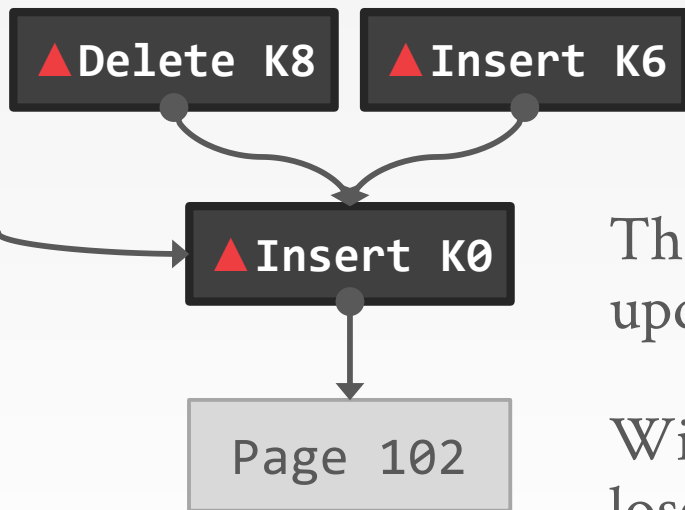
Logical
Pointer ---→

Physical
Pointer →

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	



Threads may try to install updates to same page.

Winner succeeds, any losers must retry or abort


Logical
Pointer ----->

Physical
Pointer ————>

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	



Threads may try to install updates to same page.

Winner succeeds, any losers must retry or abort

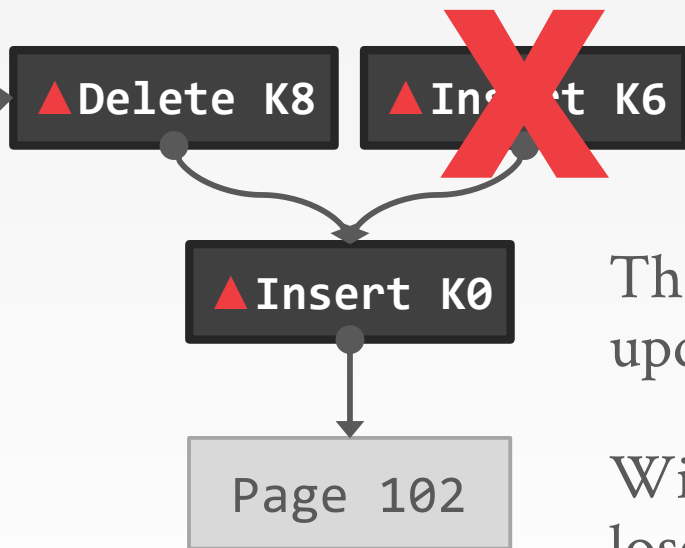
Logical
Pointer ----->

Physical
Pointer ----->

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	



Threads may try to install updates to same page.

Winner succeeds, any losers must retry or abort

Logical Pointer ----->

Physical Pointer ————>

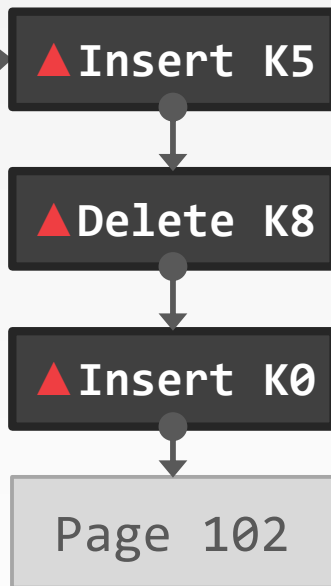
BW-TREE: CONSOLIDATION

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer 

Physical
Pointer 



Consolidate updates by creating new page with deltas applied.

BW-TREE: CONSOLIDATION

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ----->

▲ Insert K5

▲ Delete K8

▲ Insert K0

Page 102

Consolidate updates by creating new page with deltas applied.

New 102

▲ Insert K0

BW-TREE: CONSOLIDATION

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ----->

▲ Insert K5

▲ Delete K8

▲ Insert K0

Page 102

New 102

Consolidate updates by creating new page with deltas applied.

CaS-ing the mapping table address ensures no deltas are missed.

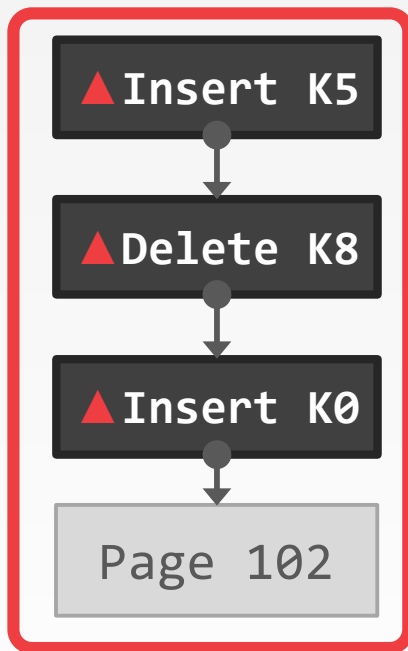
BW-TREE: CONSOLIDATION

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ----->



Consolidate updates by creating new page with deltas applied.


CaS-ing the mapping table address ensures no deltas are missed.

Old page + deltas are marked as garbage.

GARBAGE COLLECTION

We need to know when it is safe to reclaim memory for deleted nodes in a latch-free index.


- Reference Counting
- Epoch-based Reclamation
- Hazard Pointers
- Many others...



GARBAGE COLLECTION

We need to know when it is safe to reclaim memory for deleted nodes in a latch-free index.



- Reference Counting
- Epoch-based Reclamation
- Hazard Pointers
- Many others...



GARBAGE COLLECTION

We need to know when it is safe to reclaim memory for deleted nodes in a latch-free index.


- Reference Counting
- Epoch-based Reclamation
- Hazard Pointers
- Many others...



GARBAGE COLLECTION

We need to know when it is safe to reclaim memory for deleted nodes in a latch-free index.

- Reference Counting
- Epoch-based Reclamation
- Hazard Pointers
- Many others...



REFERENCE COUNTING

Maintain a counter for each node to keep track of the number of threads that are accessing it.

- Increment the counter before accessing.
- Decrement it when finished.
- A node is only safe to delete when the count is zero.

This has bad performance for multi-core CPUs

- Incrementing/decrementing counters causes a lot of cache coherence traffic.

OBSERVATION

We don't care about the actual value of the reference counter. We only need to know when it reaches zero.

We don't have to perform garbage collection immediately when the counter reaches zero.

EPOCH GARBAGE COLLECTION

Maintain a global **epoch** counter that is periodically updated (e.g., every 10 ms).

→ Keep track of what threads enter the index during an epoch and when they leave.

Mark the current epoch of a node when it is marked for deletion.

→ The node can be reclaimed once all threads have left that epoch (and all preceding epochs).

Also known as ***Read-Copy-Update*** (RCU) in Linux.

BW-TREE: GARBAGE COLLECTION

Operations are tagged with an **epoch**

- Each epoch tracks the threads that are part of it and the objects that can be reclaimed.
- Thread joins an epoch prior to each operation and post objects that can be reclaimed for the current epoch (not necessarily the one it joined)

Garbage for an epoch reclaimed only when all threads have exited the epoch.


BW-TREE: GARBAGE COLLECTION

Mapping Table


<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

*Logical
Pointer* ----->

*Physical
Pointer* ----->



Epoch Table



BW-TREE: GARBAGE COLLECTION

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>


▲ Insert K5

▲ Delete K8

▲ Insert K0


Page 102

New 102



CPU1

Epoch Table



BW-TREE: GARBAGE COLLECTION


Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>

▲ Insert K5




▲ Delete K8



▲ Insert K0

Page 102

New 102



Epoch Table



BW-TREE: GARBAGE COLLECTION


Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>

▲ Insert K5




▲ Delete K8



▲ Insert K0

Page 102

New 102



Epoch Table




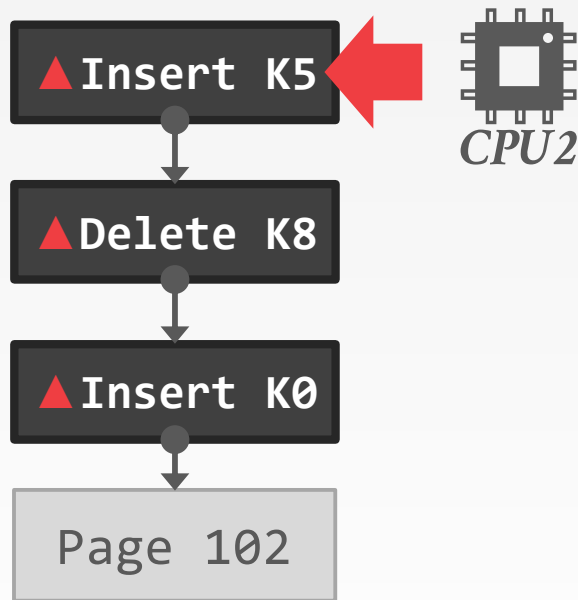
BW-TREE: GARBAGE COLLECTION

Mapping Table


PID	Addr
101	
102	
103	
104	

Logical
Pointer

Physical
Pointer



Epoch Table



BW-TREE: GARBAGE COLLECTION

Mapping Table


PID	Addr
101	
102	●
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>

New 102

▲ Insert K5




▲ Delete K8

▲ Insert K0

Page 102

Epoch Table



▲ Insert K5


▲ Delete K8

▲ Insert K0

Page 102

BW-TREE: GARBAGE COLLECTION

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer 

Physical
Pointer 


New 102

▲ Insert K5


▲ Delete K8

▲ Insert K0

Page 102



Epoch Table



▲ Insert K5

▲ Delete K8

▲ Insert K0

Page 102

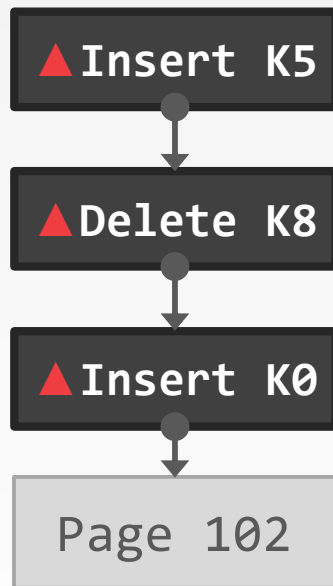
BW-TREE: GARBAGE COLLECTION

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>



New 102

Epoch Table



BW-TREE: GARBAGE COLLECTION

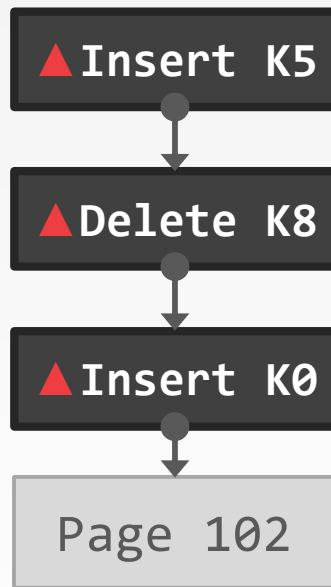
Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ---→

Physical
Pointer →

New 102



Epoch Table



BW-TREE: STRUCTURE MODIFICATIONS

Split Delta Record

- Mark that a subset of the base page's key range is now located at another page.
- Use a logical pointer to the new page.

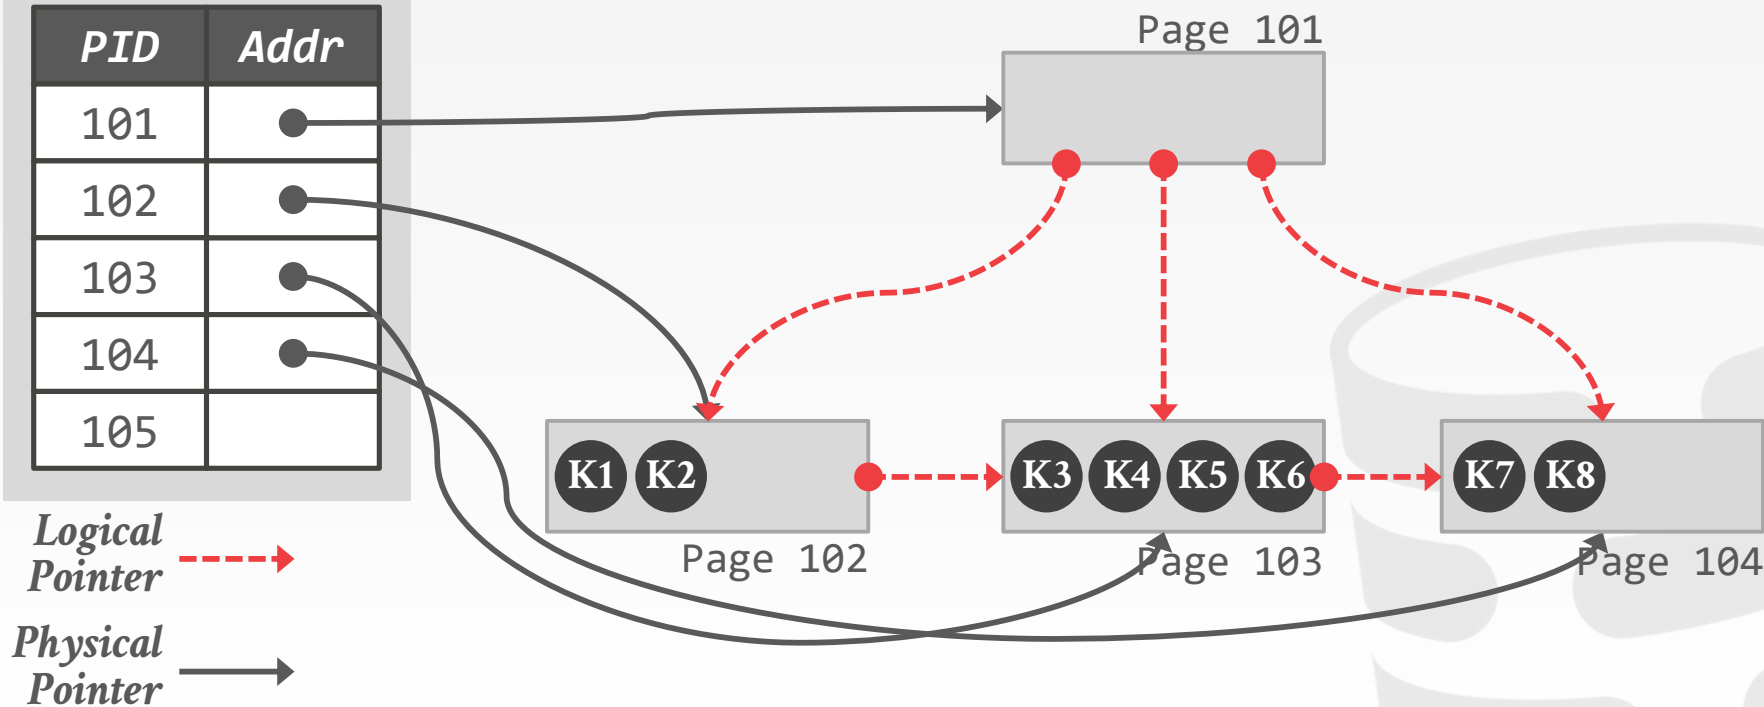
Separator Delta Record

- Provide a shortcut in the modified page's parent on what ranges to find the new page.

BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

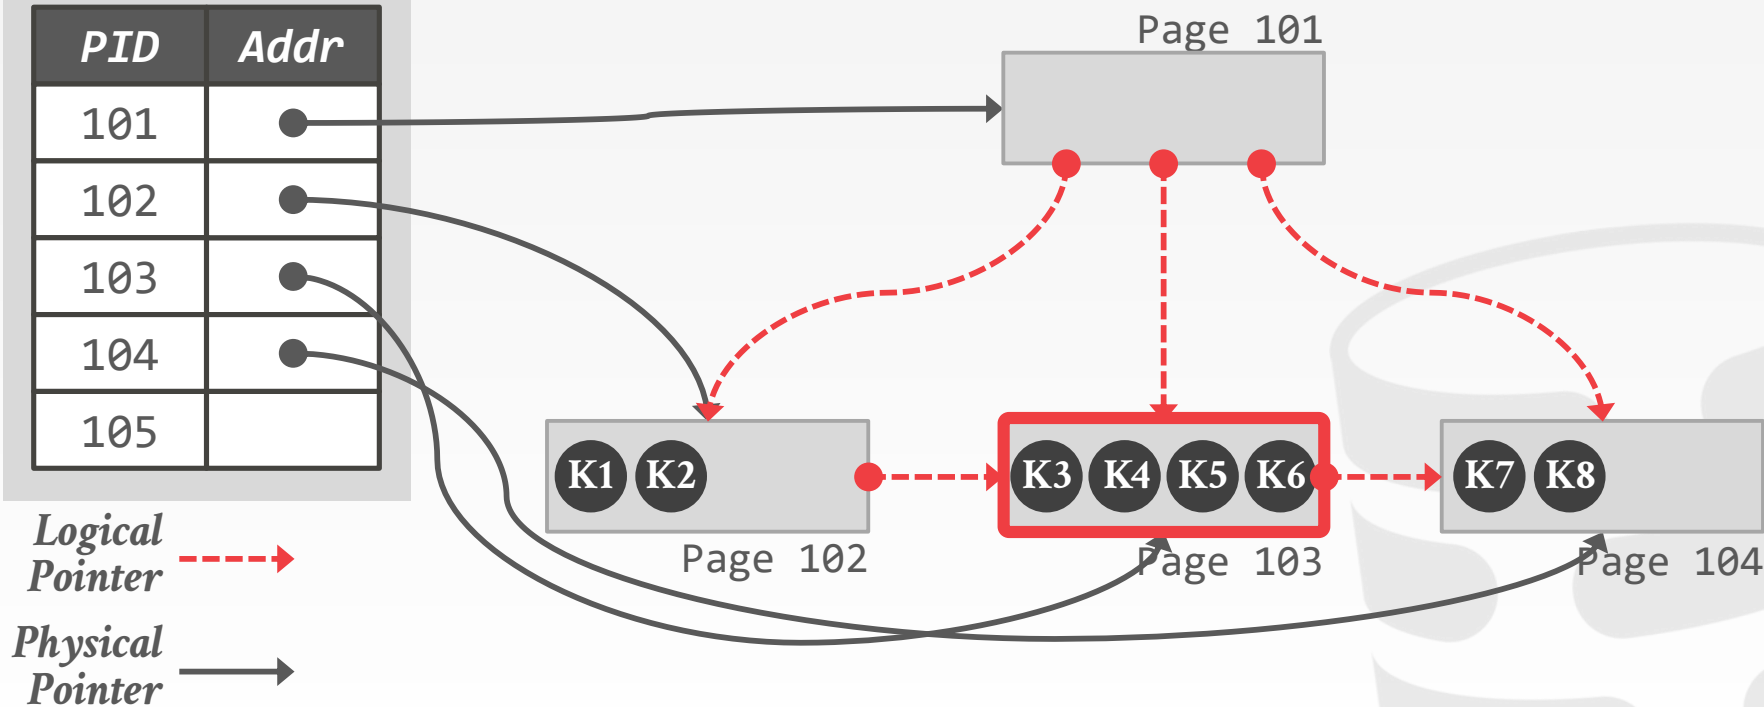
<i>PID</i>	<i>Addr</i>
101	●
102	●
103	●
104	●
105	



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

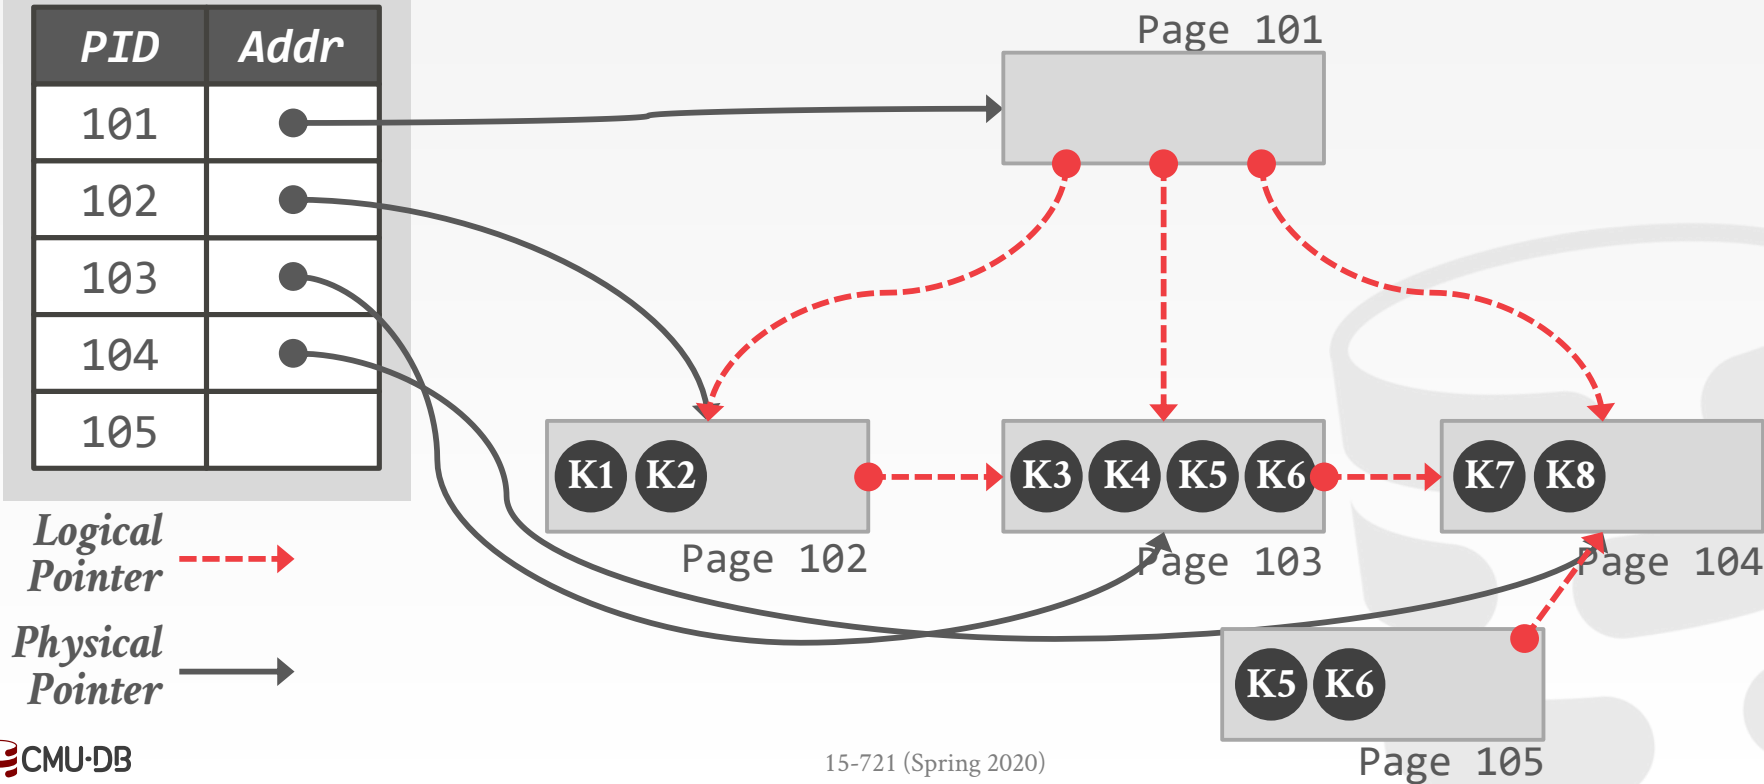
PID	Addr
101	●
102	●
103	●
104	●
105	



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

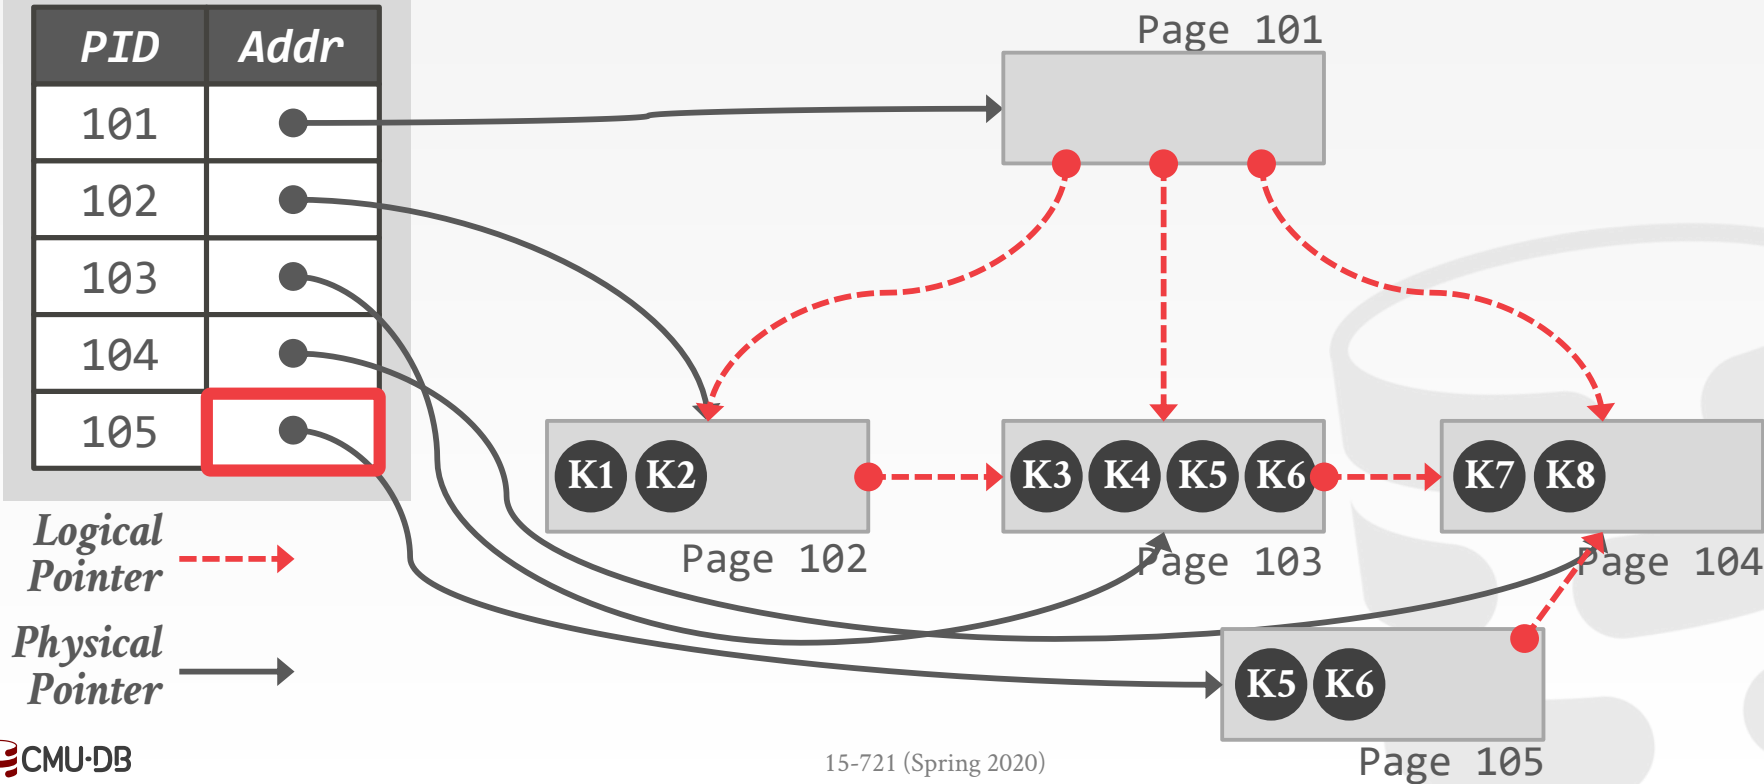
PID	Addr
101	●
102	●
103	●
104	●
105	



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

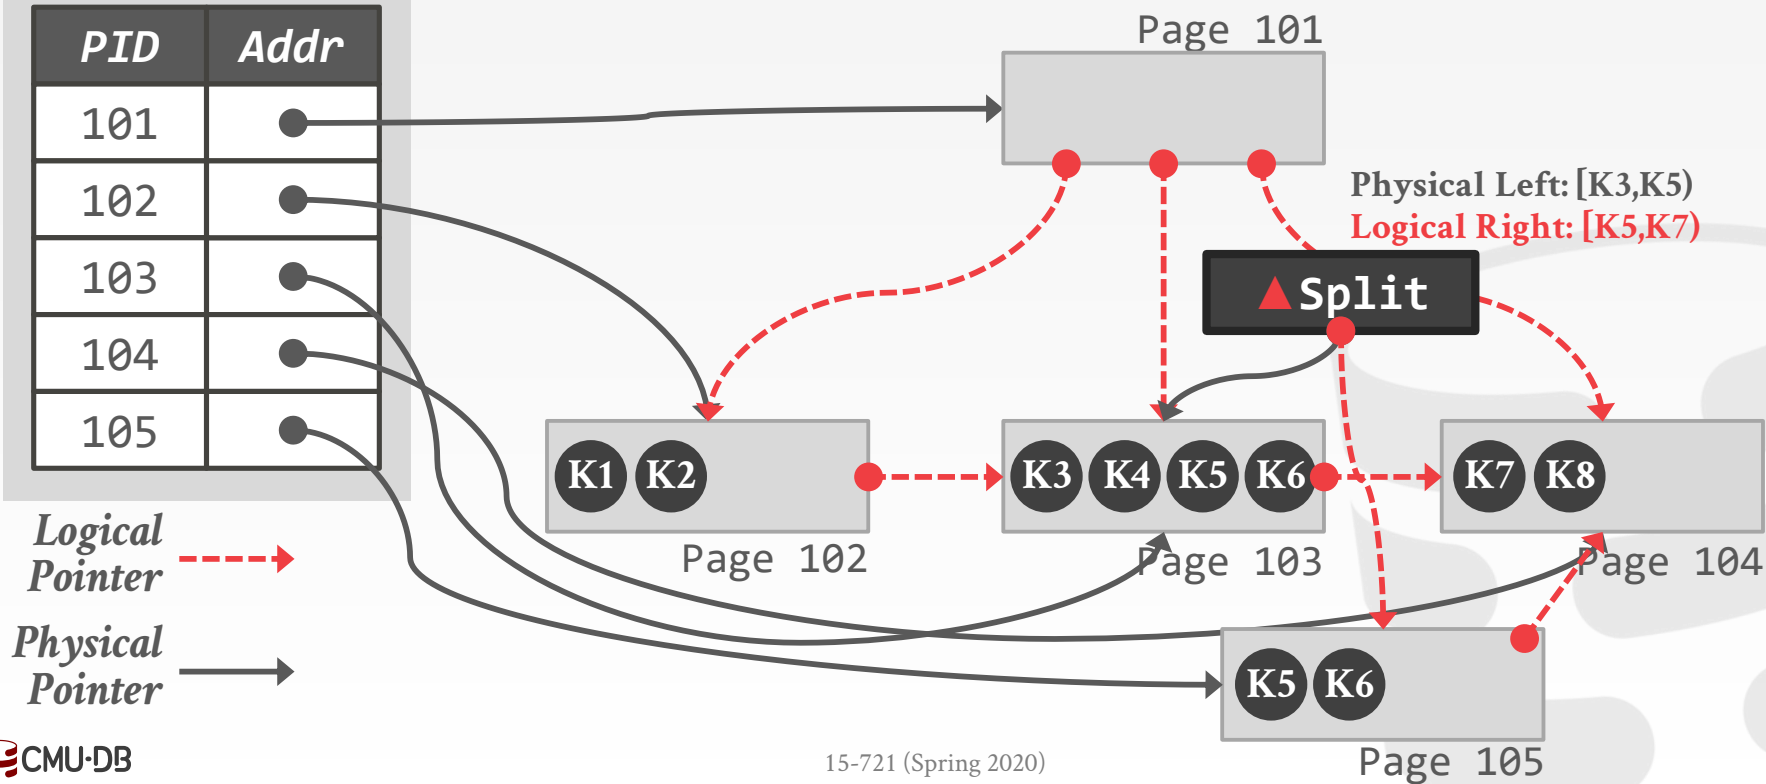
PID	Addr
101	●
102	●
103	●
104	●
105	●



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

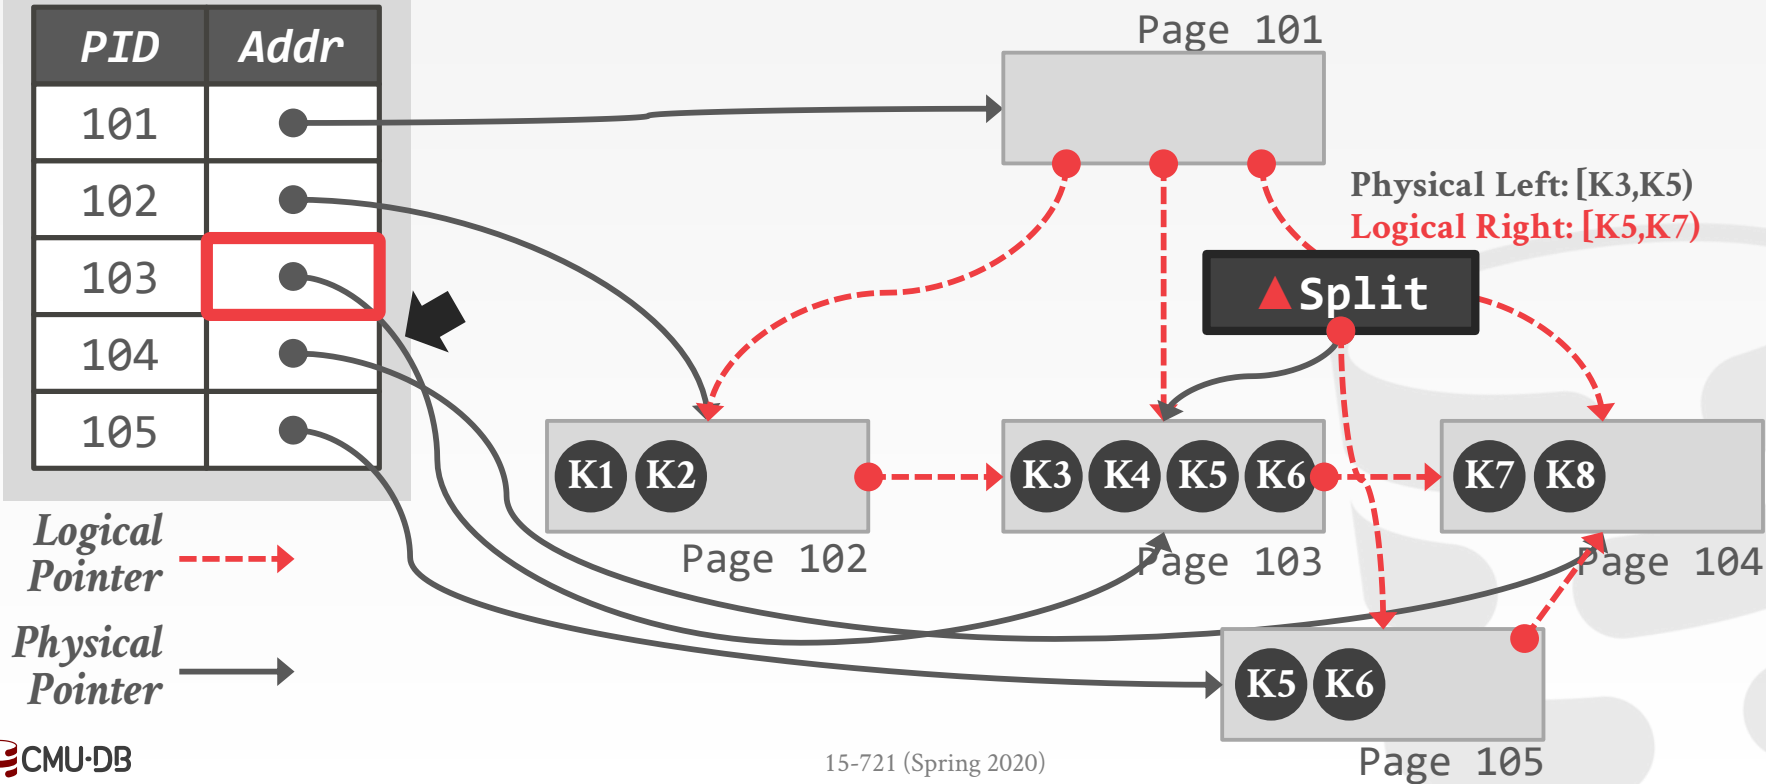
PID	Addr
101	●
102	●
103	●
104	●
105	●



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

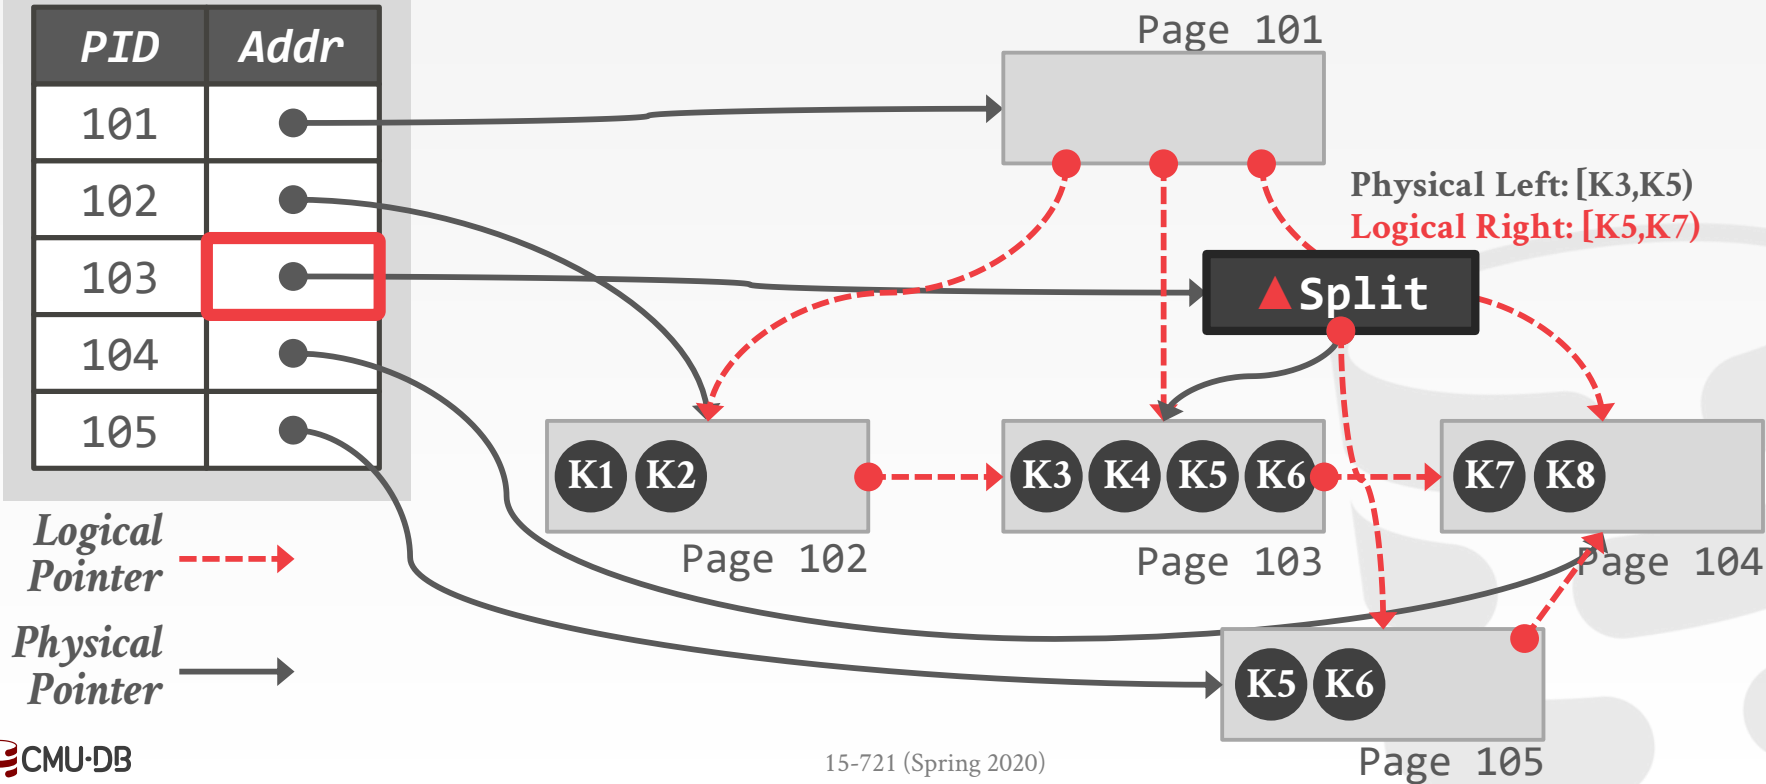
PID	Addr
101	●
102	●
103	●
104	●
105	●



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

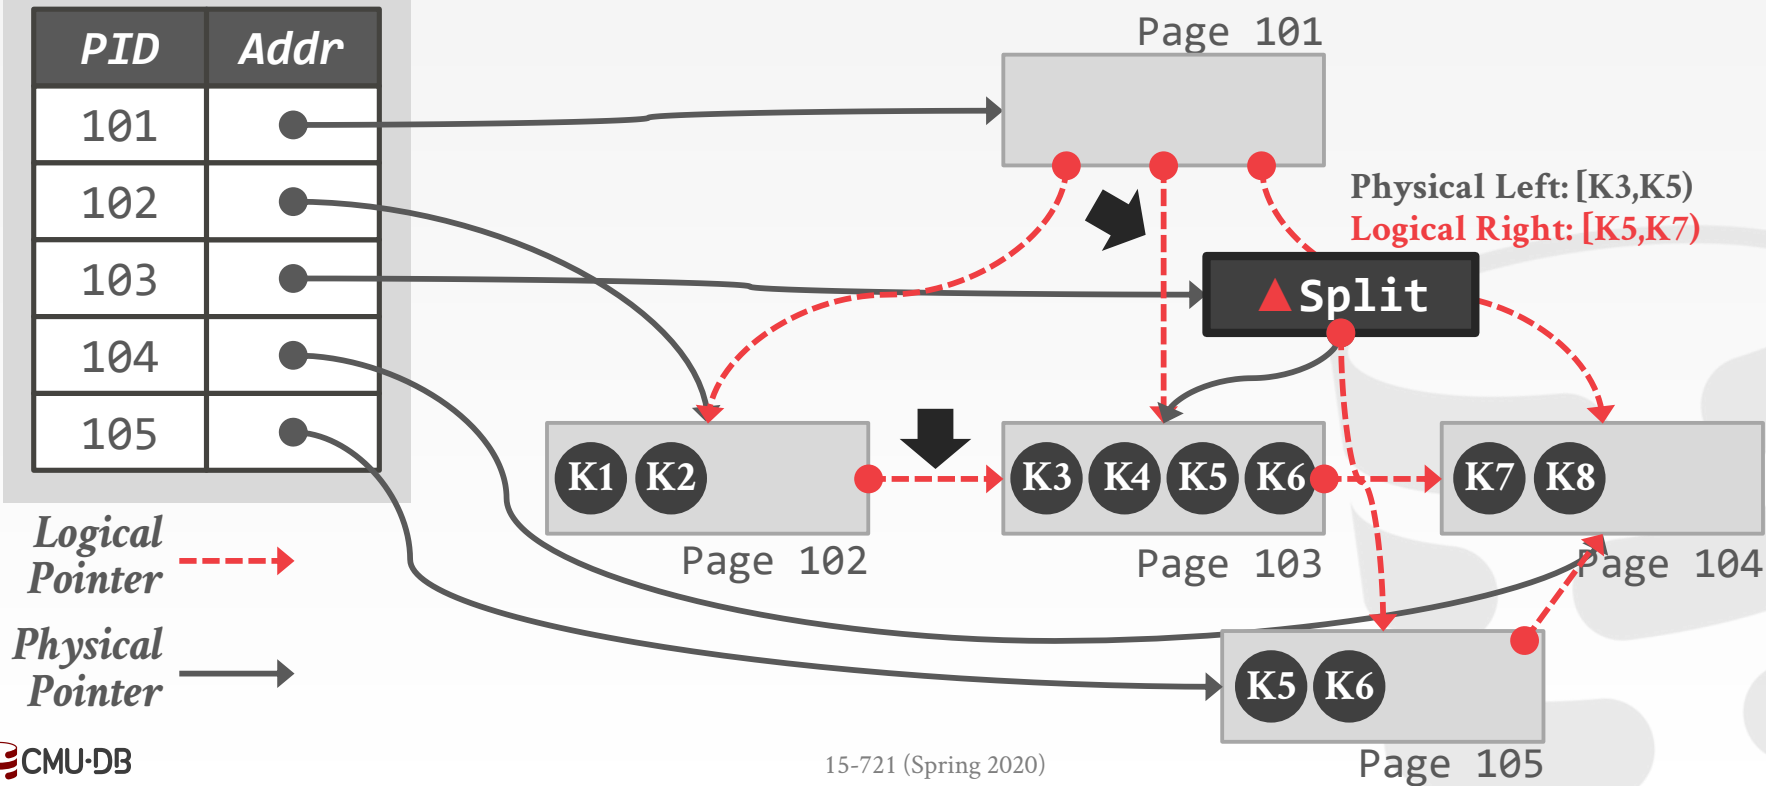
PID	Addr
101	●
102	●
103	●
104	●
105	●



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

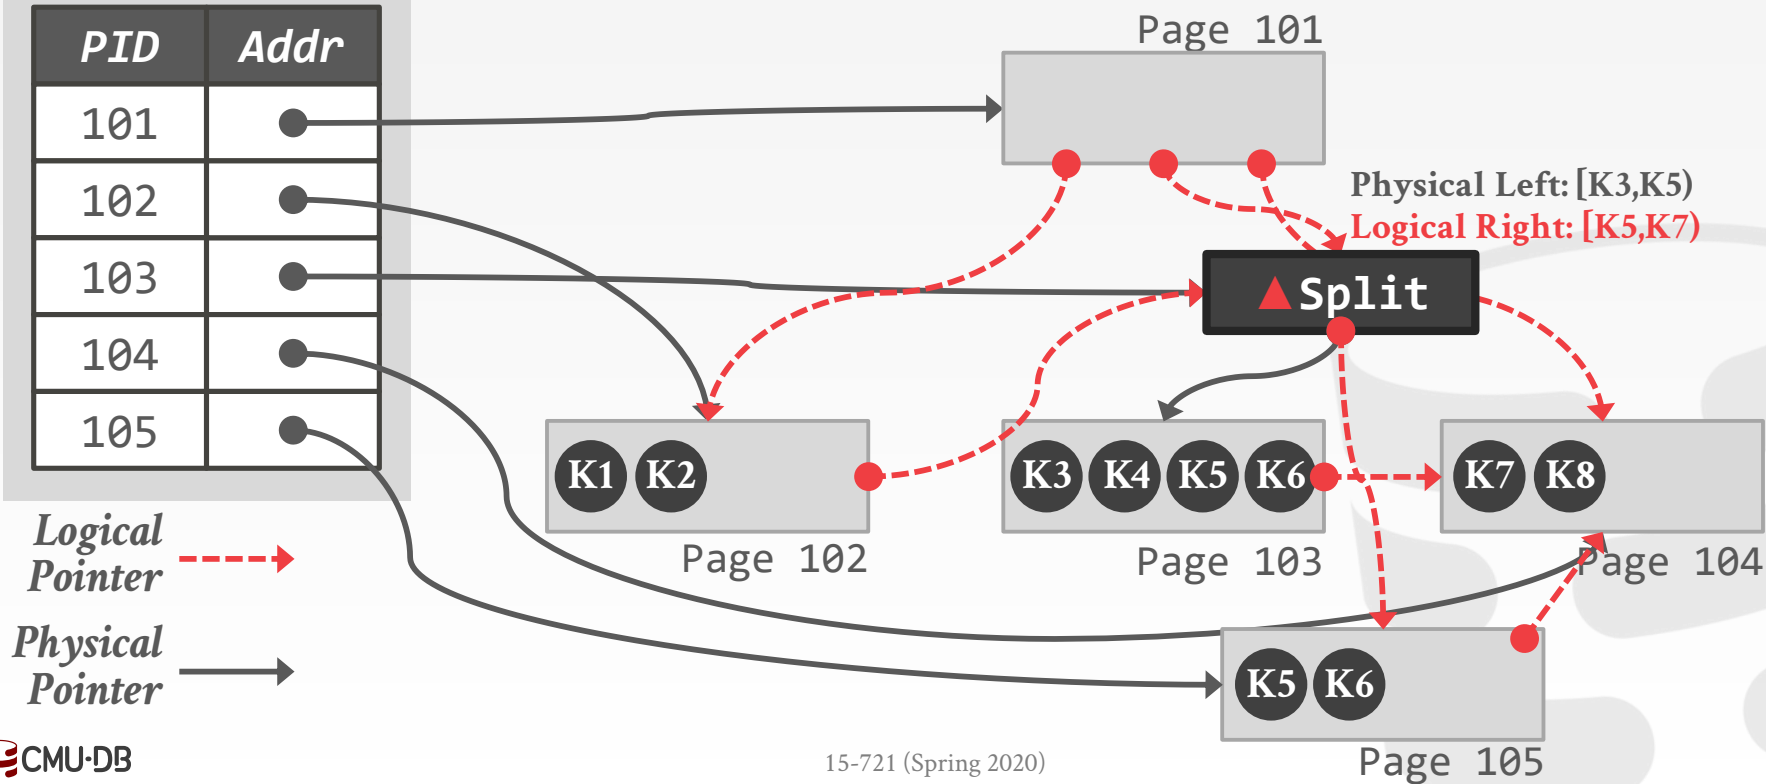
PID	Addr
101	●
102	●
103	●
104	●
105	●



BW-TREE: STRUCTURE MODIFICATIONS






Mapping Table

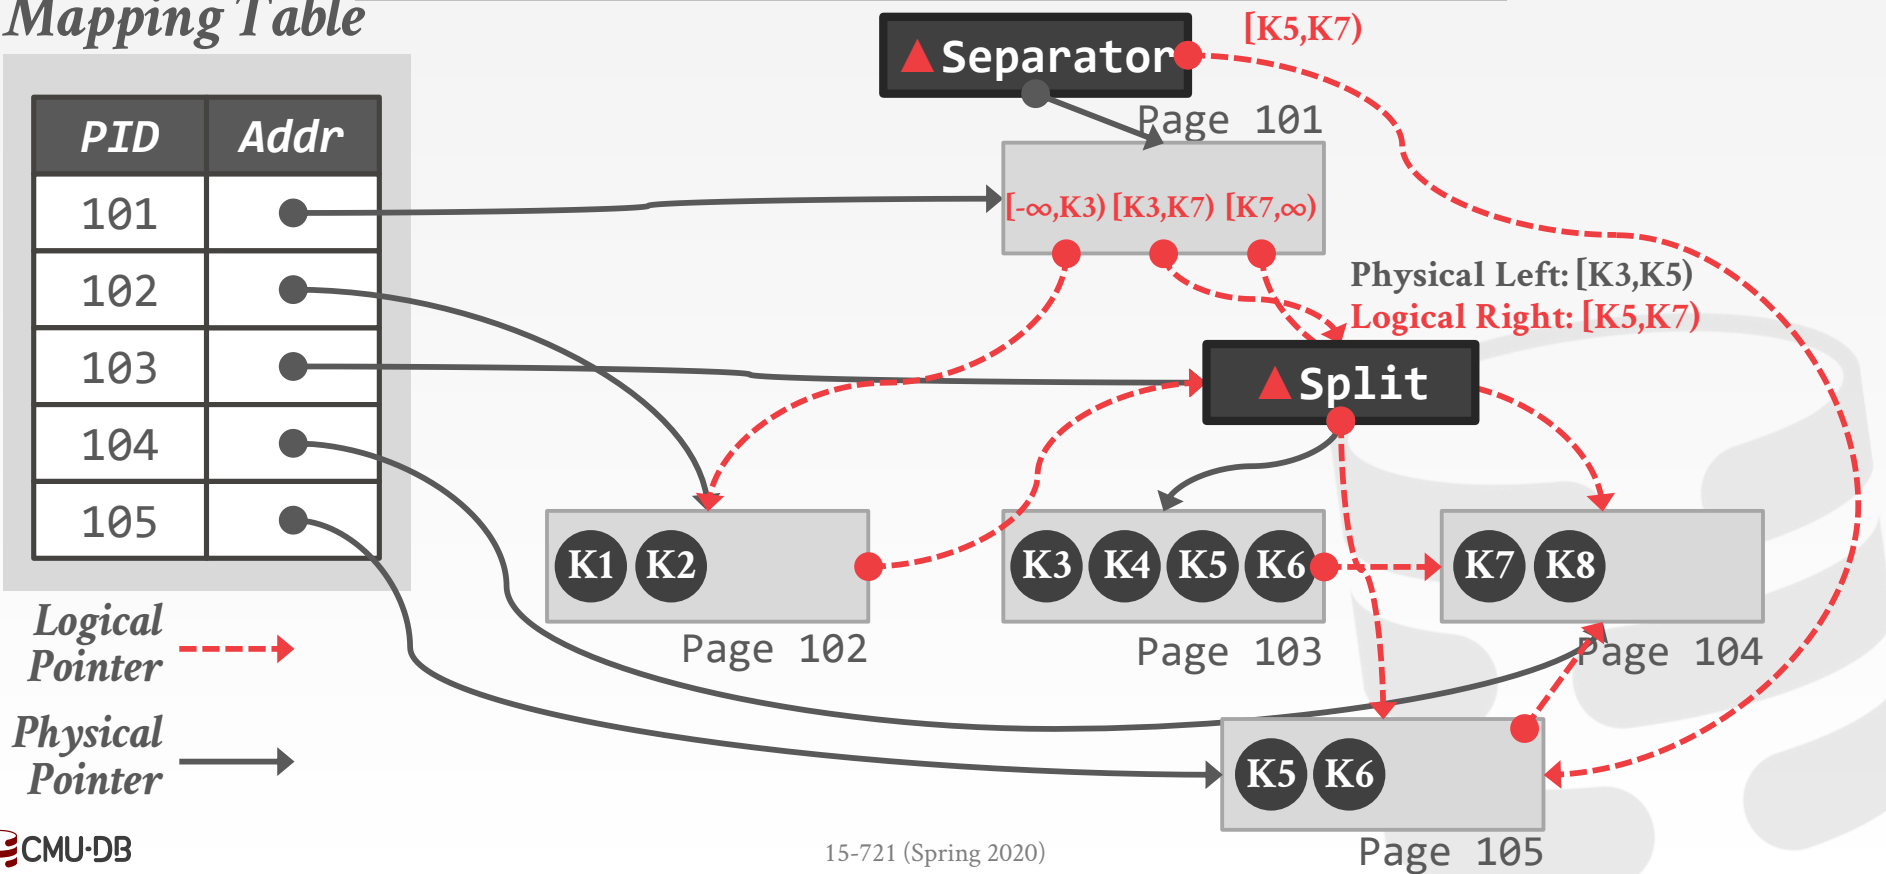
PID	Addr
101	●
102	●
103	●
104	●
105	●



BW-TREE: STRUCTURE MODIFICATIONS






Mapping Table

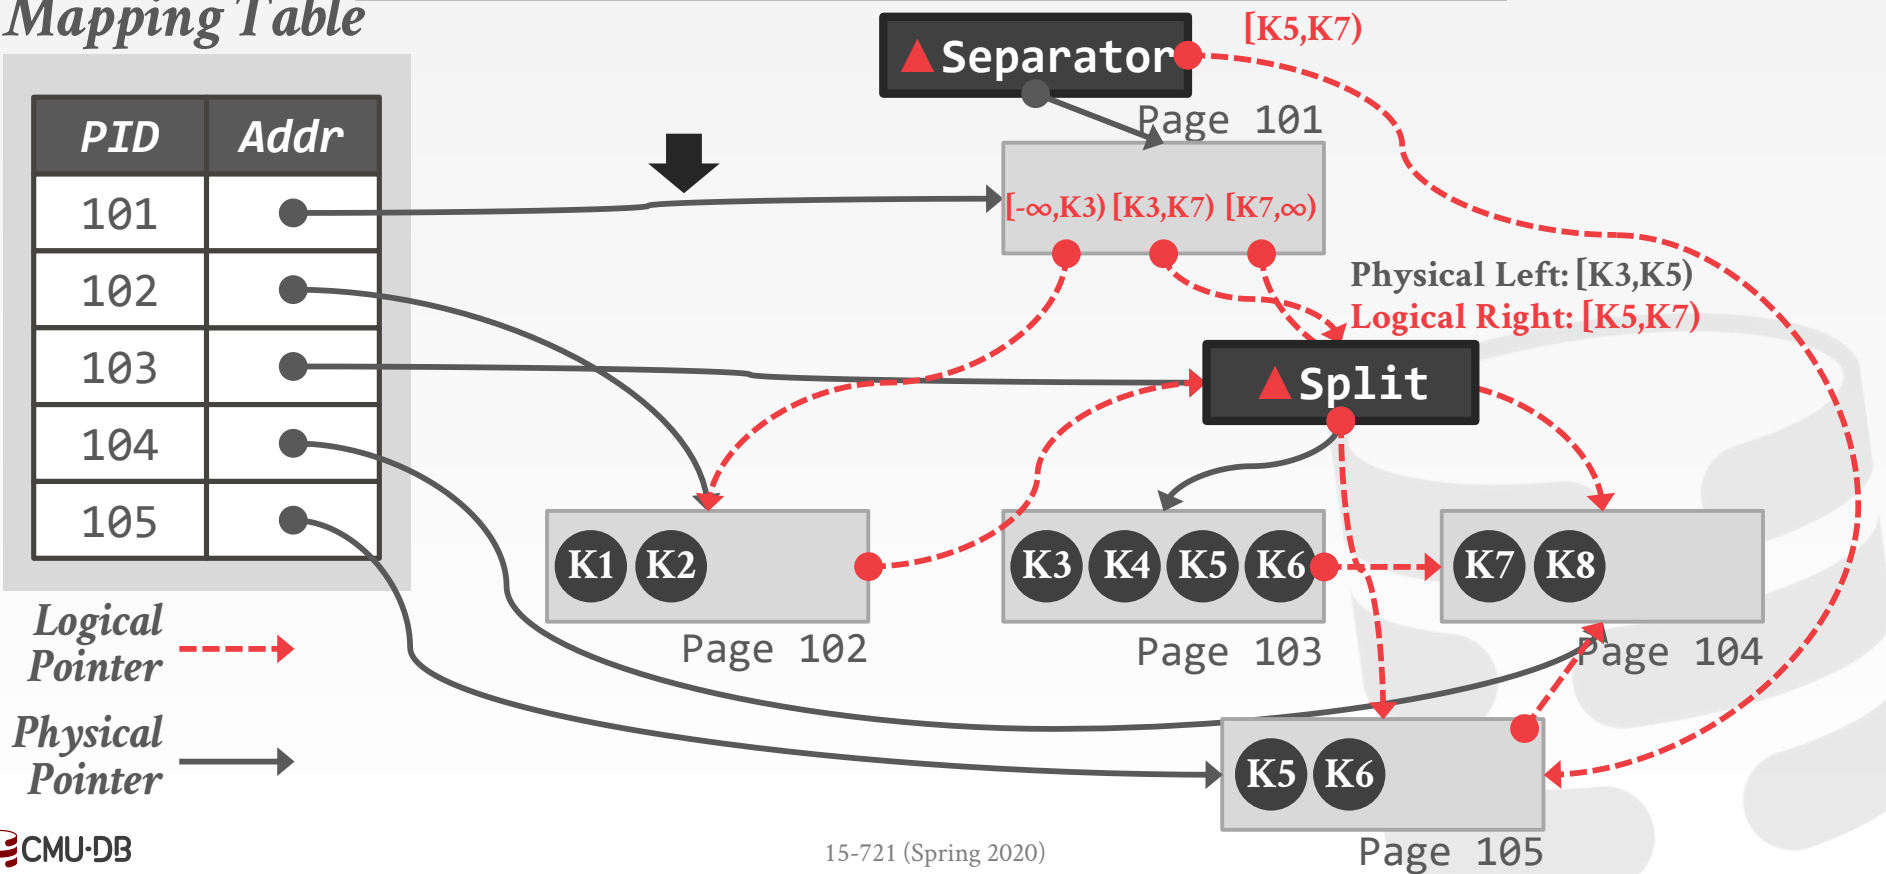
<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	
105	



BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	
105	



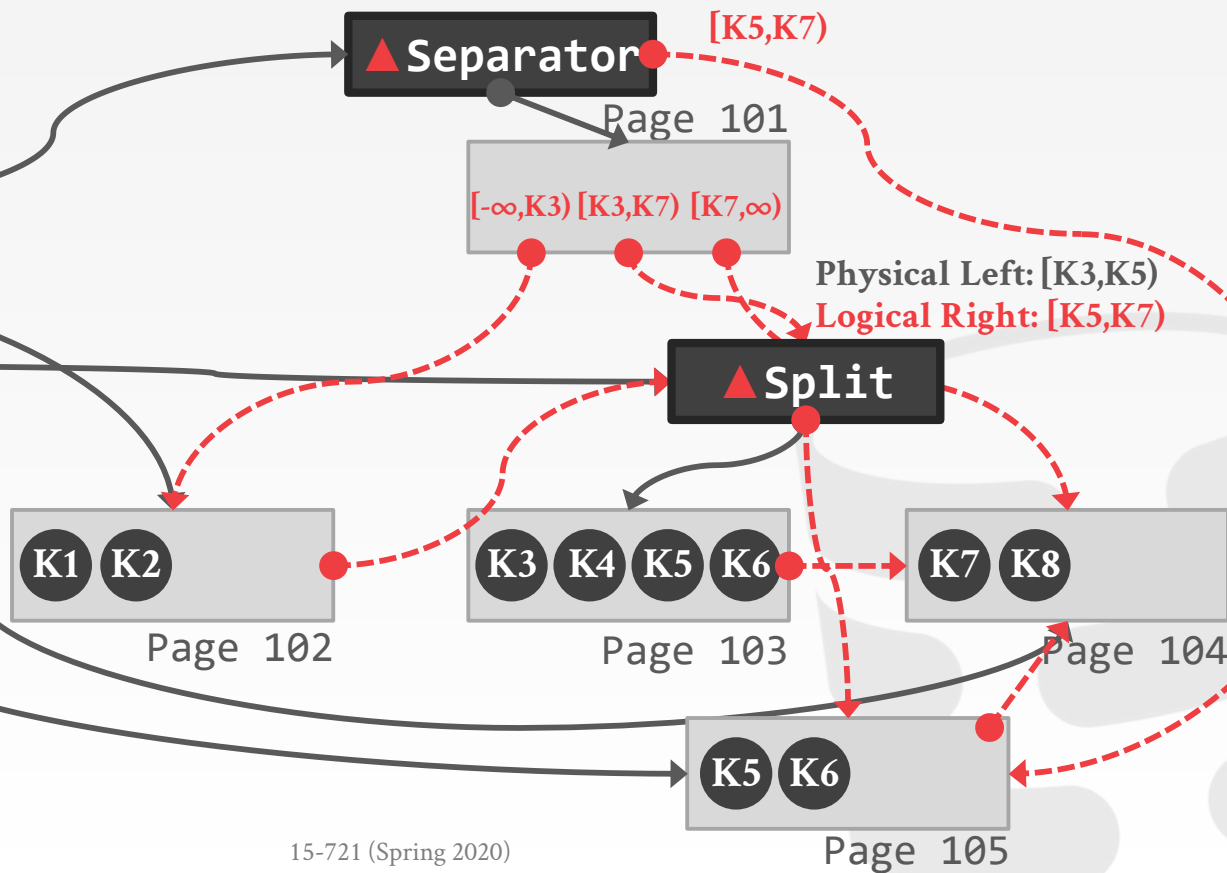
BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical
Pointer ---→

Physical
Pointer —→



CMU OPEN BW-TREE

Optimization #1: Pre-Allocated Delta Records

- Store the delta chain directly inside of a page.
- Avoids small object allocation, list traversal.

Mapping Table

<i>PID</i>	<i>Addr</i>
102	



Delta Slots


CMU OPEN BW-TREE

Optimization #1: Pre-Allocated Delta Records

- Store the delta chain directly inside of a page.
- Avoids small object allocation, list traversal.

Mapping Table

<i>PID</i>	<i>Addr</i>
102	



CMU OPEN BW-TREE

Optimization #1: Pre-Allocated Delta Records

- Store the delta chain directly inside of a page.
- Avoids small object allocation, list traversal.

Mapping Table

PID	Addr
102	●

*Delta Slot
Offset*



Delta Slots

CMU OPEN BW-TREE

Optimization #1: Pre-Allocated Delta Records

- Store the delta chain directly inside of a page.
- Avoids small object allocation, list traversal.

Mapping Table

PID	Addr
102	●

*Delta Slot
Offset*




Delta Slots

CMU OPEN BW-TREE

Optimization #1: Pre-Allocated Delta Records

- Store the delta chain directly inside of a page.
- Avoids small object allocation, list traversal.

Mapping Table

PID	Addr
102	

*Delta Slot
Offset*



Delta Slots

CMU OPEN BW-TREE

Optimization #1: Pre-Allocated Delta Records

- Store the delta chain directly inside of a page.
- Avoids small object allocation, list traversal.

Mapping Table



CMU OPEN BW-TREE

Optimization #2: Mapping Table Expansion


- Fastest associative data structure is a plain array.
- Allocating the full array for each index is wasteful
- Old Peloton: 1m nodes per index = 8MB

Use virtual memory to allocate the entire array without backing it with physical memory.

- Only need to allocate physical memory when threads access higher offsets in the array.

BW-TREE: PERFORMANCE

Processor: 1 socket, 4 cores w/ 2xHT



Source: [Justin Levandoski](#)

BW-TREE: PERFORMANCE

Processor: 1 socket, 10 cores w/ 2×HT

Workload: 50m Random Integer Keys (64-bit)




Source: [Ziqi Wang](#)

BW-TREE: PERFORMANCE

Processor: 1 socket, 10 cores w/ 2×HT

Workload: 50m Random Integer Keys (64-bit)

■ Open Bw-Tree ■ Skip List ■ B+Tree ■ Masstree ■ ART



Source: [Ziqi Wang](#)

PARTING THOUGHTS

Managing a concurrent index looks a lot like managing a database.

A Bw-Tree is hard to implement.

Versioned latch coupling provides some the benefits of optimistic methods with wasting too much work.

NEXT CLASS

Latch Implementations
Trie Data Structures

